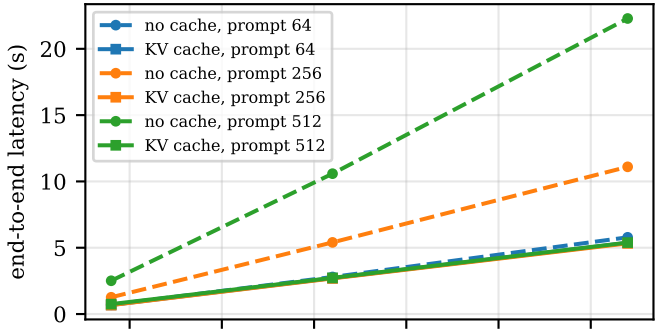
no cache, prompt 64
KV cache, prompt 64
no cache, prompt 256
KV cache, prompt 256
no cache, prompt 512
KV cache, prompt 512
end-to-end latency (s)
0
5
10
15
20
20
40
60
80
100
120
generated tokens

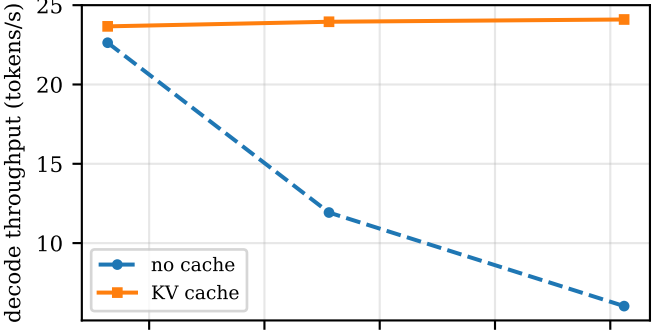
decode throughput (tokens/s)
10
15
20
25
no cache
KV cache
100
200
300
400
500
prompt length (tokens)